



Applied Bioinformatics

Short Read Alignment
October 2015, Belgrade



Goran Rakočević, PhD
goran.rakocevic@sbgenomics.com

Today's agenda

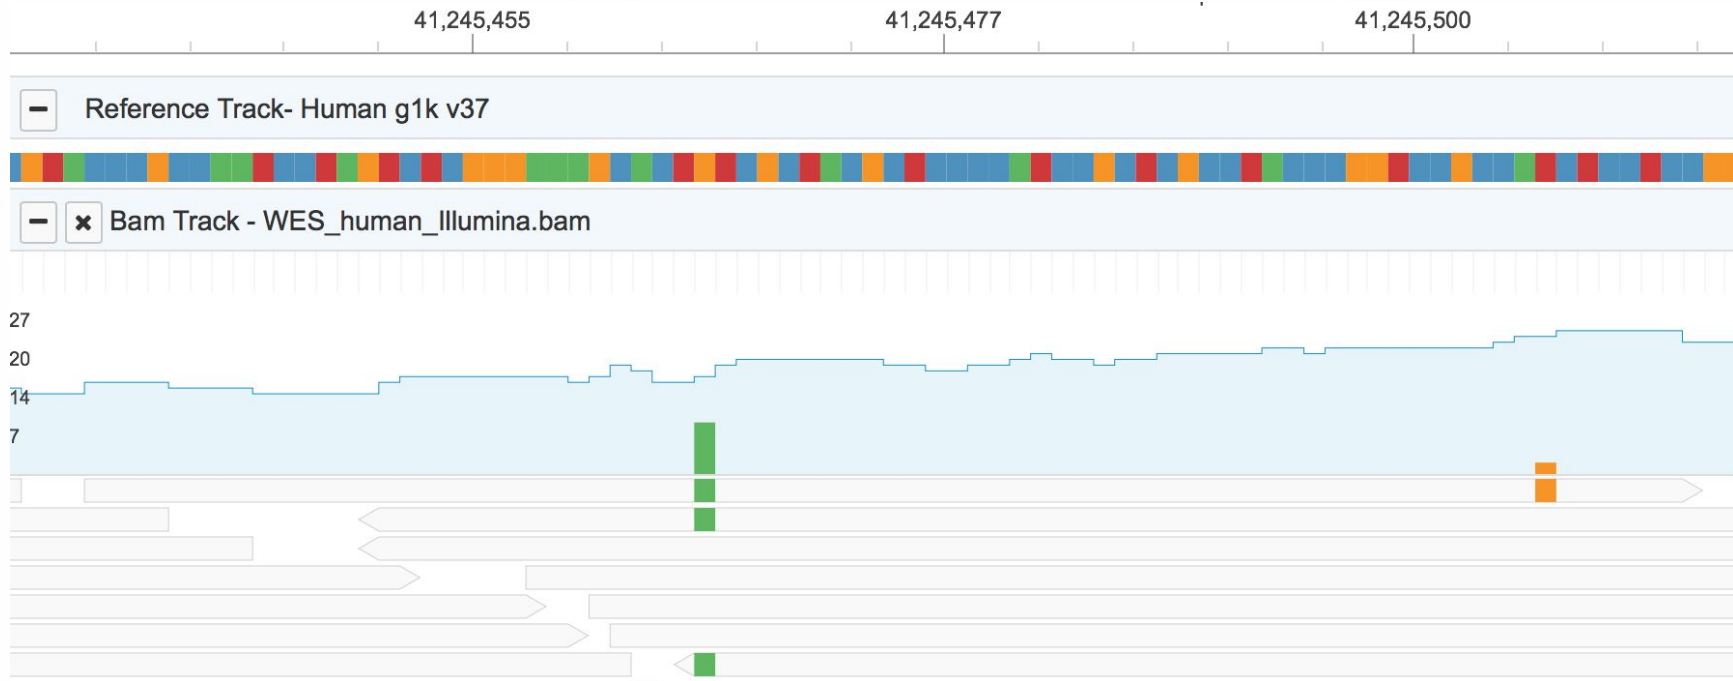
1. Overview of alignment of short reads
 - a. Two-step aligner algorithms and examples
2. Using the BWA-MEM tool
3. SAM and BAM file formats
4. Genome browsers (IGV)
5. Using PySam to explore BAM files

DNA Sequencing - Reminder

- We got a FASTQ file with the “reads” - little pieces of the genome



Aligning reads to the reference



Penalty (score) -based alignment

- Some of the reads will match perfectly to the reference
- Many reads will not:
 - Genomic variations ('mutations', SNP, Indel, etc.)
 - Sequencing errors
- Aligners commonly take a score-based approach:
 - Calculate a score, related to the distance between the read and the local reference sequence
 - Place the read so that the score is maximised
- Many algorithms exist, there is a speed/precision tradeoff

Two-step aligners

- Finding precise alignments can be expensive (for a human genome ~1 billion reads)
- Most modern aligners take a two-step approach (also called “seed and extend”)
- First find “coarse” alignments or seeds
- Then do fine grain alignments in the vicinity of the seeds
- Choose the best scoring fine grain alignment

Coarse alignment step (seeding)

- Find a set of possible coordinates
- Many false-positives, but very is usually fast
- Several common approaches:
 - K-mer hashing-based approaches
 - Radix-tree based approaches
 - FM index-based approaches (Burrows-Wheeler)
- Most require an index-building step
- A common tradeoff is speed vs RAM footprint

A simple coarse aligner

- Example: a simple hash-based scheme
- Create a hash-table which holds the positions, where each kmer in the genome occurs
- For each kmer in the read find the list of locations
- Regions with hits from many different kmers are hits
- Some kmers will have no hits, some many
- Indices quickly grow too large for whole genomes

Fine-grain alignment step (extending)

- Calculates a precise sequences match score
- Algorithms slower than core grain aligners
- Often use a lot of RAM (related to sequence size)
- The extend step also needs to produce the information on *how* the sequences match
- Mostly based on dynamic programming

Smith-Waterman aligner

- Dynamic programming local alignment algorithm

	-	P	E
-	0	0	0
E	0	0	1

Diagram illustrating the Smith-Waterman alignment algorithm. The table shows the alignment of the sequence '-' with the sequence 'P E'. The value 1 is calculated for the alignment of 'E' with 'E' (0 + 1 = 1). Arrows indicate the path from the top-left cell (0) to the bottom-right cell (1) via the top-right cell (0). The calculation 0-2=-2[0] is shown for the transition from the top-right cell (0) to the bottom-right cell (1).

$$\text{Value}_{i,j} = \begin{cases} 0 \\ \text{Value}_{i-1,j-1} + M \\ \text{Value}_{i-1,j} + G \\ \text{Value}_{i,j-1} + G \end{cases}$$

M = Match score if letters match, otherwise *mismatch penalty* (+/- 1 in this example)

G = Gap penalty (-2 in this example)

Smith-Waterman aligner

	-	P	E	R	I	C	A	A
-	0	0	0	0	0	0	0	0
E	0	0	1	0	0	0	0	0
R	0	0	0	2	0	0	0	0
C	0	0	0	0	1	1	0	0
A	0	0	0	0	0	0	2	1
A	0	0	0	0	0	0	1	3

- Calculate scores for each field
- Remember the path to each field
- Backtrack from maximum to zero
- Diagonal step is match/mismatch
- Horizontal step is a deletion
- Vertical step in an insertion

Match=1, Mismatch=-1, Gap =-2

P	E	R	I	C	A	A
-	E	R	-	C	A	A

Smith-Waterman aligner

	-	P	E	R	I	C	A	A
-	0	0	0	0	0	0	0	0
E	0	0	3	0	0	0	0	0
R	0	0	0	6	1	0	0	0
C	0	0	0	1	5	0	0	0
A	0	0	0	0	0	4	3	3
A	0	0	0	0	0	0	7	6

- The best path and the scores is defined by choice of M and G!
- Match=3, Mismatch=-1, Gap =-5

P E R I C A A
- E R C A A -

CIGAR Strings

- Run length encoding:

AAAABBBAAAD = 4A2B3A1D

- CIGAR Codes:

- **M** - alignment match (Match or Mismatch)
- **I** - insertion, **D** - Deletion, **S** - soft clip
- H, X, =, P, N - rare in DNA Seq

P E R I C A A
- E R C A A - => 1S5M1S

P E R I C A A
- E R - C A A => 1S2M1D2M

Exercise 1 - a simple aligner (30 min)

- Implement a simple hash-based aligner in Python
 - A dict can be used to create the index
- To map a read, find locations for each kmer in the read
 - Find the region with most kmers mapping to it
- A fast Smith-Waterman implementation is installed
 - Use this tool to produce fine-grain alignments
- Try modifying the to see the effect
 - Add insertions, deletions, and mismatches

BWA Aligner (BWA-MEM command)

- A widely used aligner for DNA Sequencing
- <http://bio-bwa.sourceforge.net> for more info
- BWA requires an index to be built:
 - `bwa index ref.fa`
 - Fast for small FASTA files, 2h for whole human genome
 - Produces a set of files in the same folder as the FASTA
- MEM command used for aligning
 - `bwa mem ref.fa read1.fq read2.fq > aln-pe.sam`

Exercise 2 - BWA-MEM (15 min)

- Align the example FASTQ files to the example FASTA
 - All files are in the *data* folder
- First, create the BWA index for the example FASTA file
 - Be sure to make your own copy of the FASTA!
- Align the two example paired end files using the BWA-MEM
 - Pipe the tool output to a file!
- Inspect the produced file
 - You can use less in the Jupyter notebook

BAM File format

- Standardized format for holding aligned reads
 - The read sequence and qualities
 - Position (chromosome and the first matching base)
 - CIGAR string
 - Flag (various info, like *read has a pair*, *read is aligned*, etc.)
 - Read pair position
 - Other optional tags
- Bgzip compression (roughly $\frac{1}{3}$ of raw text)
- Besides BAM, SAM (plain text) also exists

BAM File format (2)


- BAM is 0-based, SAM is 1-based
- BAM files can be sorted:
 - Coordinate sort, by read position
 - Query name (read name), read pairs are placed together
- Coordinate sorted BAM files can be *indexed*
- By convention, BAM index files:
 - Have .bai appended to the name
 - Are placed in the same folder as the BAM
 - Usually not explicitly passed to tools

Exercise 3 - SamtoBam and Sort (10 min)

- Convert the SAM file you created by BWA-MEM into BAM
 - You can use Samtools to do this
- Coordinate sort this BAM file
 - You can use Samtools or Picard to do this
- Use Samtools to view the contents of this BAM file
 - Pipe the output into less for easier viewing
- Create the bai index for this file
 - You can use Samtools or Picard to do this

The Integrative Genomics Viewer (IGV)

- IGV - A visualization tool for genomic data
- Can be used to visualize BAM, BED, and many other file types



IGV - Exercise 4 (15 min)

- Use IGV to inspect a BAM file:
 - WES_human_Illumina.bam from the *data* folder
- A BAM files needs to be sorted and indexed!
- Go to BRCA1 region and zoom to see reads
- Try different viewing options:
 - View as pairs
 - Collapsed view
 - Color by strand/insert size
- Guess where there are mutations, and what are errors!

Pysam - Python interface for BAMs

- Pysam can be used to process BAM files
- `pysam.AlignmentFile`
 - `AlignmentFile(path_to_file)`
 - `for read in AlignmentFile(path_to_file):`
- Reads are wrapped in `AlignedSegment` objects
 - `AlignedSegment` provides access to read fields and helpers
- `pysam.AlignmentFile` supports fetching regions
 - The BAM file needs to be sorted and indexed!

Pysam - Exercise 5 (30 min)

- Create an AlignmentFile object
 - Use WES_human_Illumina.bam from the *data* folder
- Take the first read from the AlignmentFile
 - Inspect the fields in the AlignedSegment
- How many unmapped reads are there in the file?
- Create a BedTool with the exome.bed from the ../data folder
- Get the regions from the BedTool that cover BRCA1 gene
- Fetch all of the reads from the BAM file mapped to BRCA1

Questions?